

## Outline

- Agents and AI
- Agent Patterns
- Applications of Agentic AI
- Multi Agent Systems
- Multi Agent System Architectures
- Agentic AI Implementation Frameworks
- LangGraph
- Agentic RAG

1 / 89

# Agentic AI and Multi Agent Systems

Dr. Shailesh Sivan

December 13, 2025

2 / 89

## Disclaimer

- Some content and images used in this presentation are sourced from the Internet and publicly available resources.
- Certain materials (text, diagrams, and examples) are generated using AI tools.
- Some content is originally created by the presenter.
- All third-party materials are used as-is from the public domain or open-access sources.
- Full credit and acknowledgment go to the respective original creators and contributors.

This presentation is intended solely for educational and academic purposes.

## Agents and AI

3 / 89

## What is Agent and Agentic AI?

## Traditional Agents vs Agentic AI Agents

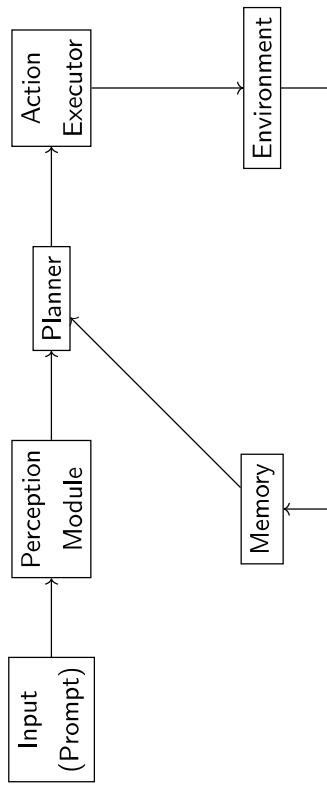
Agent	Traditional Agents	Agentic AI Agents
An <b>agent</b> is an entity that autonomously perceives its environment, makes decisions, and acts to achieve specific goals.	<ul style="list-style-type: none"><li>Rule-based / model-based intelligence</li><li>Finite-state machines, logic rules</li><li>Symbolic planning (STRIPS, PDDL)</li><li>Fixed policies and workflows</li><li>Communication protocols (ACL, FIPA)</li><li>Limited adaptability</li><li>Strong formal guarantees</li></ul>	<ul style="list-style-type: none"><li>LLM-driven reasoning and decision-making</li><li>Natural language planning</li><li>Think–Plan–Act–Reflect loops</li><li>Dynamic tool selection and composition</li><li>Language-based coordination</li><li>Self-correction via critique loops</li><li>High autonomy and generalization</li></ul>
<b>Example</b>	An AI personal assistant that schedules meetings, books travel, or completes tasks based on user instructions without being micromanaged.	Traditional agents emphasize control and correctness where as Agentic AI agents emphasize reasoning, adaptability, and autonomy.

5 / 89

6 / 89

## Key Characteristics of Agents and Agentic AI

## Simple Agent Architecture



- Goal-Directed Behavior
- Planning and Decision Making
- Autonomy and Reusability
- Memory and Adaptability
- Environment Interaction

7 / 89

8 / 89

## Types of Agents in AI

## Types of Agents in AI

- **Reflex Agents**
  - Respond only to the current percept (stateless).
  - **Example:** A thermostat that turns heating ON/OFF purely based on the current temperature reading.
- **Model-Based Agents**
  - Maintain an internal state that represents unobserved parts of the world.
  - **Example:** A robot vacuum cleaner that builds a map of the environment and navigates even in dark areas.
- **Goal-Based Agents**
  - Select actions based on the desired goal state.
  - **Example:** A navigation system (e.g., Google Maps) that plans a route to reach the destination.

### LLM-powered autonomous agent system.

- Long-term memory
- Short-term memory
- Planning
- Reflection
- Self-critics
- Chain of thoughts
- Subgoal decomposition

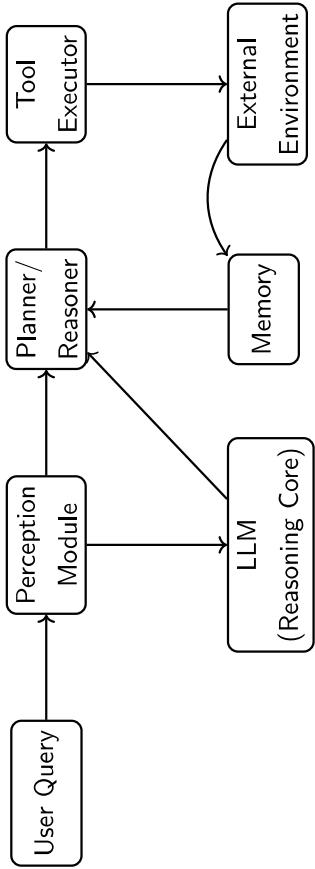
Agentic AI blends goal-based, utility-based, and learning capabilities to create autonomous and adaptive systems.

#### Note

9 / 89

10 / 89

## LLM Agent Workflow



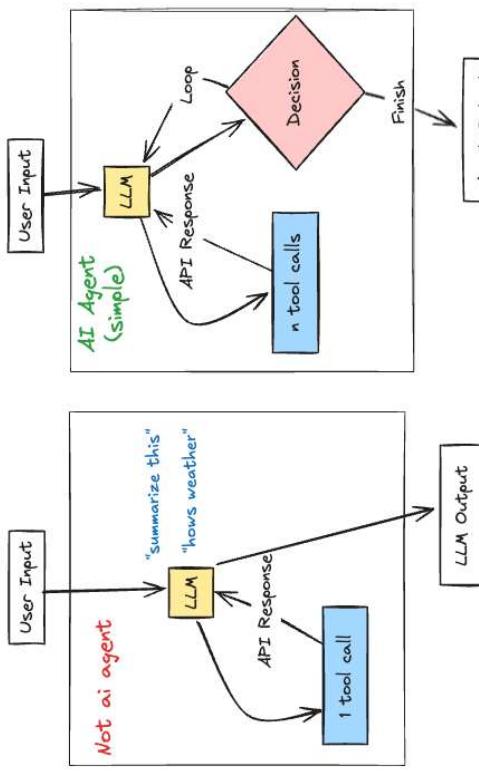
11 / 89

12 / 89

## Agent vs. Tool: Comparison

## Agent vs. Tool: Comparison

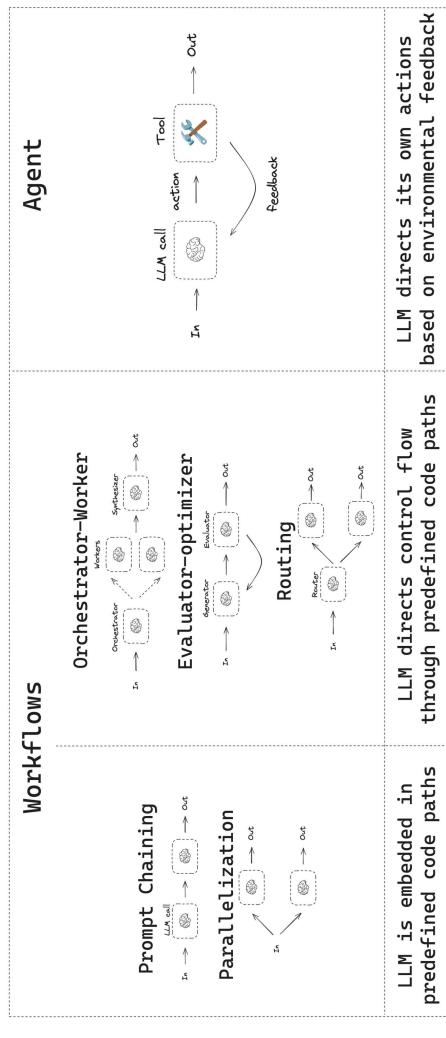
Aspect	Tool AI	AI Agent
Initiative	Reactive	Proactive
Memory	Stateless	Stateful
Autonomy	Low	High
Learning	Minimal	Continuous
Goal Orientation	Task-Specific	Goal-Driven



13 / 89

## Agents and Workflow

## AI Agents and Agentic AI



Workflows can also be implemented as Agent pattern

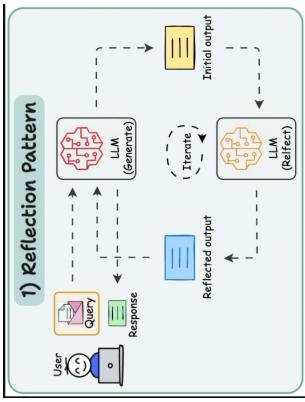
15 / 89

14 / 89

16 / 89

## Reflection Pattern

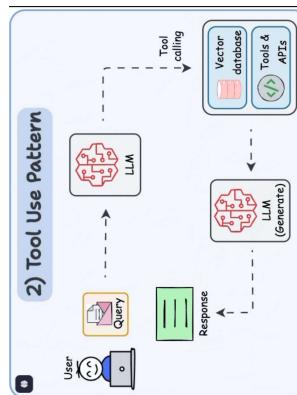
- The AI reviews its work to spot mistakes and iterate until it produces the final response.
- Self-evaluation and iterative refinement
  - Useful for hallucination reduction and answer polishing
  - Implementations: chain-of-thought + critique loops



## Agent Patterns

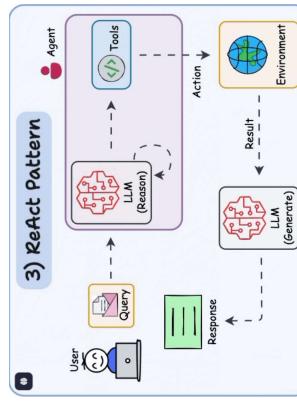
### Tool Use Pattern

- Tools allow LLMs to gather more information by:
- Querying a vector/database
  - Executing Python scripts
  - Invoking external APIs
- This reduces reliance on internal knowledge and improves factuality.



### ReAct (Reason + Act) Pattern

- ReAct combines reflection and tool use:
- Agents can reflect on generated outputs
  - Agents can call tools and observe results
  - Effective for multi-step, interactive tasks

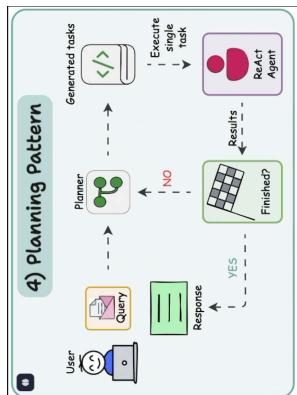


## Planning Pattern

## Multi-Agent Pattern

Instead of solving a request in one go, the AI creates a roadmap by:

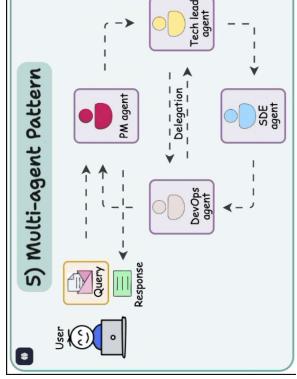
- Subdividing tasks
- Outlining objectives and steps
- Iteratively executing and revising the plan
- Strategic thinking helps solve complex tasks more effectively.



21 / 89

Multiple agents collaborate, each with dedicated roles and access to tools:

- Role-based decomposition (planner, worker, critic, router)
- Parallelization and delegation of sub-tasks
- Agents coordinate via messages or shared memory



22 / 89

## Pattern Comparison

Each pattern offers unique strengths for building intelligent systems. Depending on task complexity, environment, and goals, one or more patterns can be combined for optimal performance.

Pattern	Strengths	Best For
Reflection	Improves self-awareness and adaptability	Debugging, strategy adjustment
Tool Use	Extends capabilities via external tools	Math, retrieval, real-world tasks
Planning	Provides structure and foresight	Multi-step reasoning, code generation
Multi-Agent	Enables collaboration and modularity	Complex, distributed task environments

## Applications of Agentic AI

23 / 89

## Customer Service & Productivity

## Healthcare and Financial Services

### Customer Service

- Automated chatbots handling initial customer inquiries
- Virtual assistants providing 24/7 support
- AI agents resolving common issues without human intervention



### Healthcare

- Diagnostic support agents analyzing patient symptoms
- Medication management assistants for patients
- Administrative agents handling appointment scheduling and records



### Personal Productivity

- Digital assistants managing calendars and scheduling
- Email sorting and prioritization agents
- Task management systems with intelligent recommendations



### Financial Services

- Robo-advisors managing investment portfolios
- Fraud detection agents monitoring transactions
- Insurance claim processing automation



25 / 89

26 / 89

## Sales and Supply Chain

## Home Automation and Education

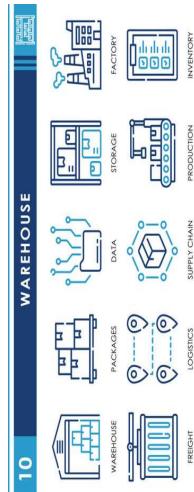
### Sales and Marketing

- Lead qualification and scoring agents
- Personalized product recommendation engines
- Marketing automation tools optimizing campaigns



### Home Automation

- Smart home assistants controlling devices and systems
- Security monitoring agents detecting unusual activity
- Energy optimization agents managing consumption



### Manufacturing and Supply Chain

- Predictive maintenance agents monitoring equipment
- Inventory optimization systems
- Quality control agents detecting defects



### Education

- Tutoring agents providing personalized learning
- Grading assistants evaluating assignments
- Course recommendation agents suggesting appropriate classes

27 / 89

28 / 89

## What is a Multi-Agent System?

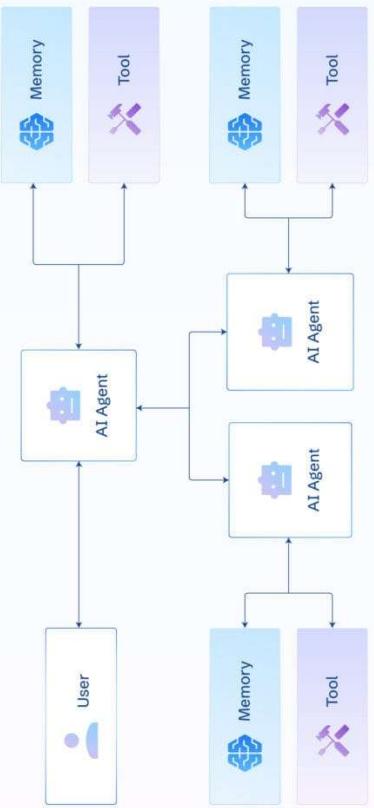
- A **multi-agent system** consists of multiple autonomous agents that collaborate or coordinate to solve a task.
- Each agent has:
  - its own goal,
  - its own reasoning logic,
  - and its own tools or knowledge.
- Agents communicate through messages and operate in a shared environment.
- Useful for:
  - decomposition of complex tasks,
  - parallel reasoning,
  - fact-checking and evaluation,
  - retrieval + reasoning + validation loops.

30 / 89

## Multi-Agent Systems

## Multi-Agent Systems

### Multi-Agent Architecture



## Single-agent vs Multi-agent

### Single-Agent Systems

- One autonomous agent interacts with environment.
- Plans, calls tools (APIs, DBs, other agents-as-tools), and responds.
- If another agent is invoked as a tool, that agent is treated as part of the environment no further modeling or coordination.

### Multi-Agent Systems

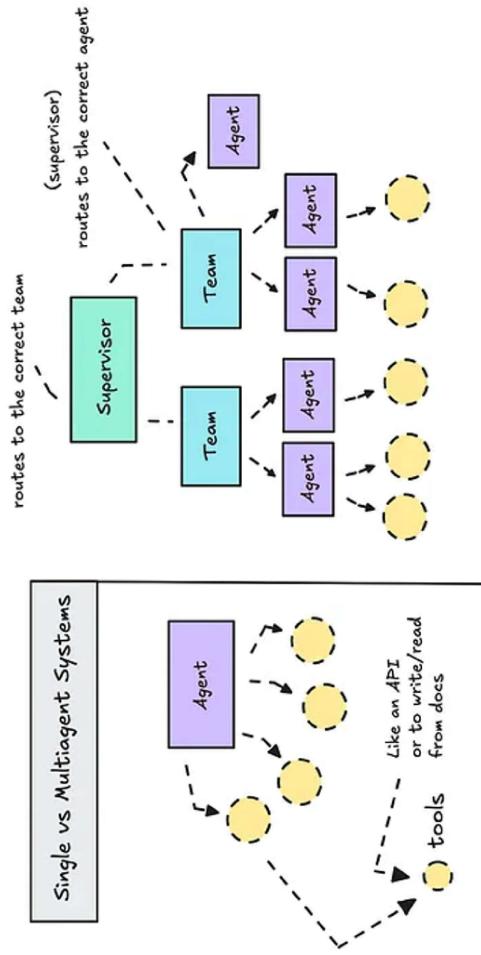
- Multiple autonomous agents that model each other's goals, memory and plans.
- Agents communicate (directly or via shared environment), cooperate and coordinate.
- Supports distributed problem solving and multi-agent reinforcement learning.

31 / 89

32 / 89

## Single-agent vs Multi-agent

## Architectures of Multi-Agent Systems



33 / 89

- **Centralized Networks**
  - A central unit stores global knowledge and orchestrates agents.
  - Pros: easy communication, uniform knowledge.
  - Cons: single point of failure.
  
- **Decentralized Networks**
  - Agents share information with neighbours; no global store.
  - Pros: robustness, modularity.
  - Cons: harder coordination and global optimization.

34 / 89

## Structures of Multi-Agent Systems

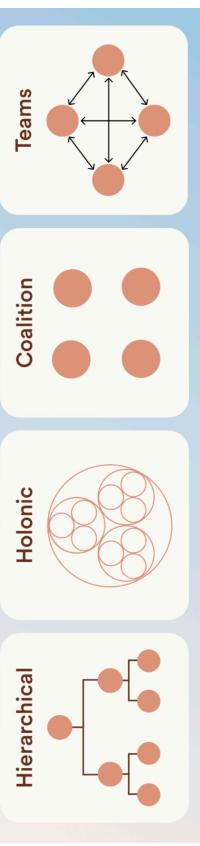
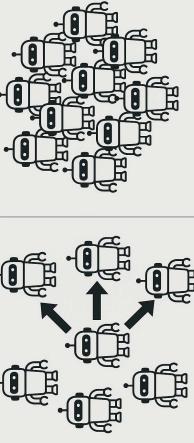
- **Hierarchical:** tree-like with varying authority levels. Can be centralized or distributed.
- **Holonic:** agents grouped into holarchies; a holon acts as a composite entity composed of subagents.
- **Coalitions:** temporary unions to boost utility; dissolve after goals met.
- **Teams:** persistent cooperative groups with interdependent roles (more dependence than coalitions).

## Agent Behaviours: Flocking & Swarming

### Flocking heuristics

- Separation: avoid collisions with nearby agents.
- Alignment: match velocity/direction of neighbours.
- Cohesion: remain close to the group.

### SWARMING



### Swarming

- Emergent self-organization with decentralized control.
- Useful when one operator manages an entire swarm (reduced operator-per-agent cost).

35 / 89

36 / 89

## Advantages of Multi-Agent Systems

- **Flexibility:** add/remove/adapt agents to changing environments.
- **Scalability:** leverage pooled information and compute.
- **Domain specialization:** different agents hold different expertise.
- **Higher performance:** collective learning, shared experiences and parallelism.

## Multi Agent System Architectures

37 / 89

### 1. Network

**Definition:** Agents can communicate directly with any other agent in the system. Each agent decides which agent to interact with next based on local state and policies.

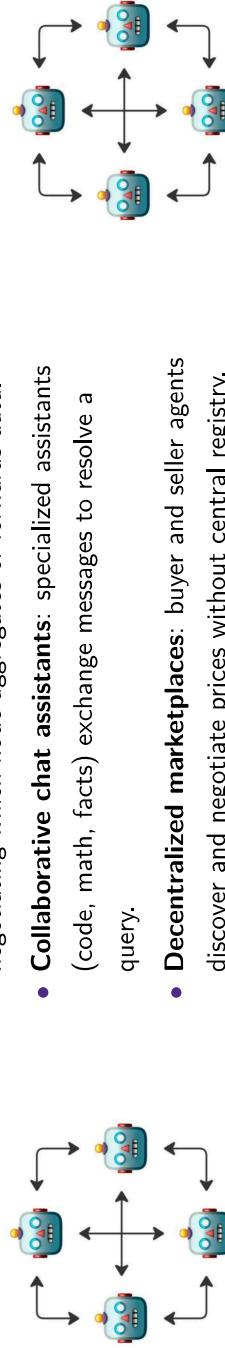
#### Key Behaviour

- Peer-to-peer communication (no central coordinator).
- Dynamic routing: agents pick interaction partners at runtime.
- Emergent cooperation via local decisions and messages.

#### Practical Use Cases

- **Ad-hoc sensor networks:** distributed sensors negotiating which node aggregates or forwards data.
- **Collaborative chat assistants:** specialized assistants (code, math, facts) exchange messages to resolve a query.
- **Decentralized marketplaces:** buyer and seller agents discover and negotiate prices without central registry.

#### Network



39 / 89

40 / 89

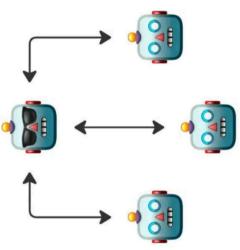
## 2. Supervisor

## 2. Supervisor

**Definition:** A central supervisor agent coordinates communication and allocates tasks among other agents. The supervisor decides which agent to call next for a given task.

### Key Behaviour

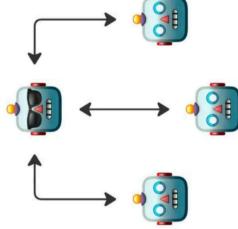
- Centralized decision-making for task routing.
- Supervisor holds or has access to global state / task queue.
- Sub-agents specialized for execution; supervisor orchestrates flow.



### Supervisor

### Practical Use Cases

- **Customer support pipeline:** supervisor assigns intents to sub-agents (billing, technical, refunds).
- **Data processing pipeline:** coordinator schedules cleaning, enrichment, analysis agents.
- **Multi-step automation:** supervisor ensures ordered execution (validate compute publish).



41 / 89

## 3. Supervisor with Tool Calling

## 3. Supervisor with Tool Calling

**Definition:** The supervisor treats individual agents (or capabilities) as tools/APIs. Instead of delegating to a live agent, the supervisor invokes tool endpoints (or agent-as-service) as needed.

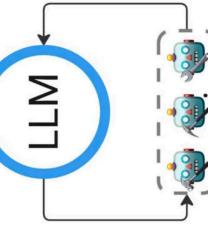
### Key Behaviour

- Supervisor orchestrates calls to tool-interfaces (search API, calculator, domain models).
- Tools are stateless or minimally stateful; logic lives in supervisor or tool.
- Easier integration with existing services and safer sandboxing.

### Supervisors (as Tools)

### Practical Use Cases

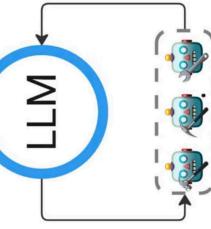
- **RAG-enabled QA:** supervisor calls vector-search, web-scraper, or DB connector tool to answer a query.



### LLM

### Supervisors (as Tools)

- **Enterprise automation:** supervisor invokes HR payroll API, CRM updates, reporting tools.
- **Robotics control:** supervisor calls motion-planning service, vision API, and safety-check tool sequentially.



42 / 89

43 / 89

44 / 89

## 4. Hierarchical

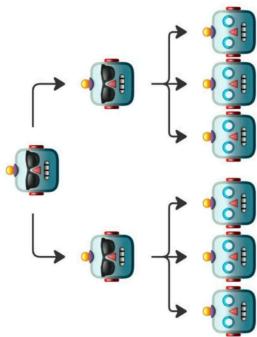
## 4. Hierarchical

**Definition:** A multi-tiered system where supervisors manage groups of agents or other supervisors/nodes in control of complex workflows and better scalability.

### Key Behaviour

- Layered supervision: top-level planners / mid-level coordinators / worker agents.
- Task decomposition and delegation across layers.
- Error isolation and logging scoped to layers.

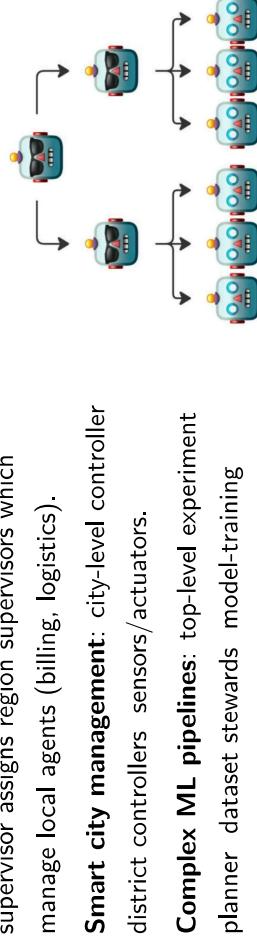
### Hierarchical



### Practical Use Cases

- **Large-scale enterprise orchestration:** global supervisor assigns region supervisors which manage local agents (billing, logistics).
- **Smart city management:** city-level controller district controllers sensors / actuators.
- **Complex ML pipelines:** top-level experiment planner dataset stewards model-training workers.

### Hierarchical



45 / 89

46 / 89

## 5. Custom Workflow

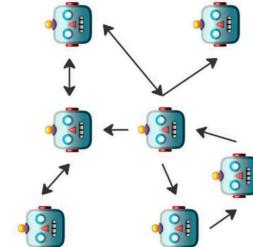
## 5. Custom Workflow

**Definition:** Agents communicate only with a subset of other agents according to predefined rules; workflows combine deterministic routing and agent decision-making.

### Key Behaviour

- Hybrid routing: deterministic rules for critical steps, agent autonomy for exploratory steps.
- Role-based visibility: agents only see subset of global state.
- Rule engine and agent policies co-exist.

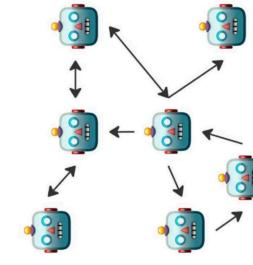
### Custom



### Practical Use Cases

- **Regulated workflows:** banking KYC deterministic verification then agent-based risk assessment.
- **Healthcare triage:** fixed triage rules route critical cases, non-critical go to specialized agents for extended assessment.
- **Hybrid manufacturing lines:** fixed safety checks plus agentic workers optimizing assembly substeps.

### Custom



47 / 89

48 / 89

## Decision Guide & Trade-offs

## Challenges of Multi-Agent Systems

### Choosing the right pattern

- **Network** choose when decentralization and robustness are required; expect harder coordination.
- **Supervisor** choose when deterministic task ordering or centralized policy is needed.
- **Supervisor + Tools** choose to integrate existing services and reduce inter-agent complexity.
- **Hierarchical** choose for scale and error isolation in large deployments.
- **Custom Workflow** choose for regulated domains where partial determinism is required.

**Common concerns:** latency, single points of failure, security of inter-agent channels, observability and debugging.

49 / 89

50 / 89

## Popular Frameworks



- LangChain + LangGraph
- AutoGPT / BabyAGI
- ReAct (Reasoning + Acting)
- CrewAI / AgentVerse

## Agentic AI Implementation Frameworks

52 / 89

Core Libraries

- Popular Languages
    - Python (dominant in LLM tools)
    - JavaScript (for front-end agents)
    - TypeScript (used in LangChain.js)
  - LangChain, LangGraph
  - AutoGPT, BabyAGI
  - CrewAI, OpenAgents
  - Transformers ( )
  - ReAct, Semantic Kernel

53 / 89

## Simple Agent (search + calc) example

```
from langchain.google_genai import ChatGoogleGenerativeAI
from langchain.agents import Tool, initialize_agent, AgentTool

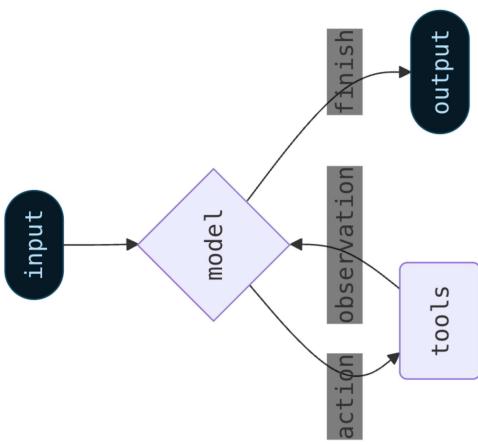
def calc_tool(expr: str) -> str:
    return str(eval(expr))

def search_tool(q: str) -> str:
    return "SEARCH_RESULT: product X released in 2010"

t1=Tool(name="search", func=search_tool, description="Search
info and return results.")

t2=Tool(name="calc", func=calc_tool, description="Evaluate
expressions safely.")
```

54 / 85



54 / 85

## Simple Agent (search + calc) example

```
llm = ChatGoogleGenerativeAI(
    model="gemini-2.0-flash", temperature=0,
    google_api_key="YOUR_GOOGLE_API_KEY"
)

agent = initialize_agent(tools=[t1, t2],
    llm=llm,
    agent=AgentType.ZERO_SHOT_REACT_DESCRIPTION,
    verbose=True
)
```

55 / 89

## Simple Python Agentic AI Example

## Agent Demo: Weather-Based Match Scheduling

### A basic autonomous agent that answers and searches

```
from langchain.google_genai import ChatGoogleGenerativeAI
from langchain.agents import initialize_agent, load_tools

llm = ChatGoogleGenerativeAI(model="gemini-2.0-flash", temperature=0,
    google_api_key="YOUR_GOOGLE_API_KEY")

tools = load_tools(["serpapi", "llm-math"], llm=llm, serpapi_api_key="YOUR_API_KEY")

agent = initialize_agent(tools=tools, llm=llm, agent="zero-shot-react-description", verbose=True)

result = agent.run("What's the weather in Paris and square root of 123?")
print(result)
```

57 / 89

## Prompt Used by the Match-Decision Agent

### Agent Role + Inputs

You are MATCH-DECIDER, an agent that decides whether to schedule an outdoor match tomorrow at {location}.  
Inputs:  
- forecast: {weather\_json}  
- constraints: thresholds for rain, temperature, wind, player preferences.

### Decision Rules + Output

Task:  
1) Decide: "PLAY" or "POSTPONE".  
2) Give 2-3 bullet reasons.  
3) If POSTPONE -> suggest alternate slot.  
Rules:  
- Rain chance >= threshold => POSTPONE.  
- Temp < min\_temp => POSTPONE.  
- Wind > wind\_limit => POSTPONE.  
Output JSON:  
{ "decision": "...", "reason": "..." }

### Purpose:

- Fetch tomorrows weather for a location
  - Return simplified JSON used by LLM
- Inputs:**
- location name
  - date (implicitly tomorrow)

### Code

```
from langchain.tools import tool
import requests, json, datetime
@tool
def weather_tool(location: str):
    lat, lon = 42.36, -71.06 # Example coords
    url = f"https://api.openweathermap.org/..."
    r = requests.get(url)
    data = r.json()
    tomorrow = (datetime.date.today() + datetime.timedelta(days=1)).isoformat()
    return json.dumps({"date": tomorrow, "forecast": data.get("daily", [])})
```

## Weather Tool (Tomorrow's Forecast)

### Goal: Build an agent that checks

tomorrows weather and decides: **PLAY** or **POSTPONE** a match.

### Agent Workflow:

- Input: location, date (tomorrow)
- Node 1: WeatherTool fetch forecast
- Node 2: LLM Reasoning evaluate conditions
  - Node 3: DecisionTool return PLAY/POSTPONE
  - Optional: Loops for self-correction
  - Agent graph

58 / 89

59 / 89

60 / 89

## Decision Tool + Assembling the Agent

## Decision Tool + Assembling the Agent

### Decision Tool

```
@tool
def scheduler_decision(forecast_json: str, constraints: str):
    fc = json.loads(forecast_json)
    cons = json.loads(constraints)
    daily = fc["forecast"][0]
    temp = daily["temp"]["min"]
    rain = daily["pop"]
    wind = daily["wind_speed"]
    if rain >= cons["precip_threshold"]:
        or temp < cons["min_temp"]
        or wind > cons["wind_limit"]:
            return {"action": "POSTPONE"}
    return {"action": "PLAY"}
```

61 / 89

### What Makes Agent Prompts Different?

### Building the Agent

```
from langchain import ChatOpenAI
from langchain.agents import create_agent

model = ChatOpenAI(model="gpt-4o-mini")
tools = [weather_tool, scheduler_decision]
agent = create_agent(model=model, tools=tools)

result = agent.run({
    "input": "Should we play tomorrow?",
    "location": "Bengaluru",
    "constraints": {
        "min_temp": 15,
        "precip_threshold": 0.3,
        "wind_limit": 8
    }
})
print(result)
```

62 / 89

### Agent Identity (Role Prompt)

Clearly define what the agent is and its level of expertise.

#### Template:

```
You are <AGENT_NAME>, a highly reliable <ROLE>.
Your purpose is to <PRIMARY_OBJECTIVE>.
```

#### Example:

```
You are DATA-CURATOR, an expert in data cleaning.
Your purpose is to clean and validate datasets for ML workflows.
```

63 / 89

64 / 89

## Goals and Success Criteria

## Tools the Agent Can Use

Agents work better with measurable objectives.

### Template:

#### Example:

Your goals:  
1. <Goal 1>  
2. <Goal 2>  
3. <Goal 3>

Success means:  
- <Success Condition 1>  
- <Success Condition 2>

Define available tools and when to use them.

### Template:

#### Example:

You can use the following tools:  
- <Tool A>: <What it does>  
- <Tool B>: <What it does>

Use tools ONLY when needed **and** return valid tool input.

### Success means:

- No missing values.
- Transformations are logged.

### Success means:

- No missing values.
- Transformations are logged.

#### Example:

Tools:  
- PythonTool: For running Python code.  
- WebSearch: For searching information online.

65 / 89

## Action Loop / Reasoning Style

Define how the agent should Think Plan Act Reflect.

### Template:

#### Follow this loop:

1. THINK: Analyze the state.
  2. PLAN: Break the goal into steps.
  3. ACT: Use a tool **or** output text.
  4. REFLECT: Evaluate **and** improve.
- Repeat until goal **is** achieved.

### Example:

If the tool result **is** wrong, reflect **and** retry.

## Constraints and Output Format

Prevent hallucinations with strict rules.

### Template:

#### Constraints:

- Do **not** fabricate facts.
  - Do **not** assume missing data.
- Output Format:
- ```
{  
  "step": "...",  
  "action": "...",  
  "result": "..."  
}
```

### Example:

Never generate synthetic data unless allowed.  
Return output **as** structured JSON.

67 / 89

68 / 89

66 / 89

## Complete Sample Agent Prompt

## Best Practices for Agent Prompt Design

A ready-to-use agent prompt.

### Role, Goals, Success

You are RESEARCH-ASSISTANT, an academic research agent.

**GOALS:**

1. Summarize literature.
2. Identify research gaps.
3. Propose methodology.

**SUCCESS:**

- All info cited.
- Output structured.

**TOOLS:**

- WebSearch
- PythonTool

### Tools, Loop, Constraints, Output

**ACTION LOOP:**

1. THINK
2. PLAN
3. ACT
4. REFLECT

**CONSTRAINTS:**

- No hallucinated citations.

**OUTPUT FORMAT:**

```
{ "summary": "...",
  "gaps": [...],
  "methodology": "..." }
```

69 / 89

70 / 89

## Why LangGraph?

- Traditional LangChain agents are powerful but:
  - hard to debug,
  - unpredictable (LLM decides planning),
  - stateless across steps,
  - not suited for multi-agent orchestration.

- **LangGraph** solves these by:

- Representing agent workflows as **graphs** (nodes = steps),
- Adding **stateful memory** across runs,
- Supporting **multi-agent collaboration**,
- Allowing **interrupt/resume**, retries, and human-in-loop.

LangGraph

72 / 89

## What is LangGraph?

## LangGraph: Basic Structure

- **State** Shared memory during execution.
- **Nodes** Computation units (LLM calls, tools, retrievers).
- **Edges** Execution transitions.
- **Graph** Connect nodes to define an agent workflow.

### Example workflow:

1. User input
2. Retrieval node (FAISS)
3. Reasoning node (Gemini)
4. Output node

73 / 89

## Sample LangGraph Implementation

```
import os
from pydantic import BaseModel, Field
from langgraph.graph import StateGraph, END
from langchain_google_genai import ChatGoogleGenerativeAI

# ----- State -----
class State(BaseModel):
    input: str = Field(...)
    output: str | None = None
# ----- Gemini LLM -----
llm = ChatGoogleGenerativeAI(
    model="gemini-2.0-flash",
    temperature=0,
    google_api_key=YOUR_GEMINI_API_KEY
)
```

74 / 89

## Sample LangGraph Implementation

```
# ----- Node -----
def respond(state: State) -> State:
    raw = llm.invoke(state.input)
    text = raw.content if hasattr(raw, "content") else str(raw)
    return State(input=state.input, output=text)

# ----- Build Graph -----
graph = StateGraph(State)
graph.add_node("respond", respond)
graph.set_entry_point("respond")
graph.add_edge("respond", END)
app = graph.compile()

result = app.invoke({"input": "Explain LangGraph in one line"})
print(result['output'])
```

76 / 89

## Agentic RAG Flow in LangGraph

- Node 1: **Plan Step** Decide: retrieve or compute?
- Node 2: **Retriever Step** Retrieve passages via FAISS
- Node 3: **Reason Step (Gemini)** Synthesize and answer
- Node 4: **Verification Step** Critic agent checks correctness

Agentic RAG

Each node is explicit and predictable unlike classical agents.

78 / 89

## Multi-Agent Architecture (LangGraph + Gemini)

- **Planner Agent** Determines if retrieval is required and dispatches actions.
- **Retriever Agent** Uses FAISS to fetch evidence from local documents.
- **Reasoner Agent (Gemini)** Synthesizes an answer using LLM reasoning.
- **Critic Agent** Evaluates quality, correctness, and completeness.
- LangGraph coordinates these agents through a **stateful graph**.

## Multi-Agent Workflow using LangGraph + Gemini

```
class MASState(BaseModel):  
    query: str  
    docs: list | None = None  
    draft: str | None = None  
    final: str | None = None  
    action: str | None = None  
  
    # --- Planner Agent ---  
    def planner(state):  
        if "year" in state.query or "released" in state.query:  
            state.action = "retrieve"  
        else:  
            state.action = "reason_direct"  
  
    return state
```

79 / 89

80 / 89

## Multi-Agent Workflow using LangGraph + Gemini

## Multi-Agent Workflow using LangGraph + Gemini + Gemini

```
# --- Retriever Agent ---
def retriever_agent(state):
    state.docs = db.similarity_search(state.query)
    return state

# --- Reasoner Agent (Gemini) ---
def reasoner_agent(state):
    ctx = "\n".join(d.page_content for d in (state.docs or []))
    prompt = f"Context: {ctx}\nQuestion: {state.query}\nAnswer: "
    raw = llm.invoke(prompt)
    state.draft = str(getattr(raw, "content", raw))
    return state
```

81 / 89

```
# --- Critic Agent ---
def critic_agent(state):
    critique = llm.invoke(
        f"Evaluate this answer:\n{state.draft}\n\nProvide corrections only."
    )

    state.final = str(getattr(critique, "content", critique))

    return state

# --- Build Graph ---
graph = StateGraph(MASSState)
graph.add_node("planner", planner)
graph.add_node("retriever", retriever_agent)
graph.add_node("reasoner", reasoner_agent)
graph.add_node("critic", critic_agent)
```

82 / 89

## Multi-Agent Workflow using LangGraph + Gemini

## Running the Multi-Agent System

```
result = app.invoke({"query": "When was Product X released?"})
print("Draft answer:", result.draft)
print("Final corrected answer:", result.final)
```

**Pipeline Summary:**

- Planner decides retrieval required.
- Retriever fetches documents from FAISS.
- Reasoner Gemini generates answer.
- Critic Gemini evaluates and improves.

```
graph.set_entry_point("planner")
graph.add_conditional_edges("planner",
    lambda s: s.action,
    {"retrieve": "retriever", "reason_direct": "reasoner"})
graph.add_edge("retriever", "reasoner")
graph.add_edge("reasoner", "critic")
graph.add_edge("critic", END)
app = graph.compile()
```

83 / 89

84 / 89

## Why Multi-Agent Systems in LangGraph?

## Benefits of Multi-Agent Systems in LangGraph

- LangGraph allows you to model agents as **graph nodes**.
- Each node can:
  - run different LLMs (e.g., Gemini),
  - use different tools or vector stores,
  - maintain separate state.
- Edges define:
  - communication paths,
  - collaboration rules,
  - control flow.
- Perfect for:
  - Planner → Retriever → Solver,
  - Solver → Critic → Refiner loops,
  - multi-turn workflows with checkpoints.

85 / 89

## Challenges

## Conclusion

- **LangGraph** provides a reliable and structured framework for building advanced agentic systems.
- **Multi-agent workflows** (Planner, Retriever, Reasoner, Critic) improve accuracy, robustness, and interpretability.
- **Gemini LLMs** combine high-quality reasoning with efficient tool use, making them ideal for multi-agent and RAG-based pipelines.
- **FAISS integration** enables fast, scalable retrieval to ground answers in factual context.
- Overall: LangGraph + Gemini = a powerful foundation for **autonomous, verifiable, and production-ready Agentic AI**.

86 / 89

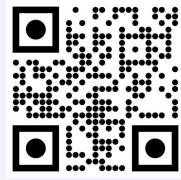
87 / 89

88 / 89

# Thank You!

*For your attention and participation !!!*

**Dr. Shailesh Sivan**



**Email:** shaileshsivan@gmail.com

**Phone:** +91-8907230664

**Website:** www.shaileshsivan.info